

Meeting Outline:

- Progress Update
 - Results:
 - SMC steering - not BART, T5
 - method
 - Twist method, which scales particles' weights with a factor at each timestep (higher is more probable) to indicate whether the particle is on the right track. Undoes the scaling after each timestep.
 - scaled by probability of target bias predicted by the bias model
 - llama 2
 - Ben running
 - gpt2
 - Done
 - gpt-neo
 - Ellie running
- Final report
 - Finish by Sunday night
 - Goal - 6/8 pages total for final paper - ~2-3 more pages for progress - (this submission should be ~5 pgs.)
 - Introduction (1.5 columns)
 - elaborate a bit more
 - preview the different levels of intervention
 - Methodology (2 pages)
 - Results
 - figure/table of the results
 - can put placeholders for the progress report
 - half page
 - Rebecca Lin
 - Write & interpret results - (results & discussion section)
 - Ben
 - Discussion
 - limitations of political bias classifier and how it might affect results
 - @reinaldo
 - Factual consistency & relevance of summaries is important
 - future work is to include this factor in the study
 - Ellie
 - Compute limitations
 - making SMC steering only generate up to 3 sentences could make it easier to steer so that's a confounding factor
 - Ellie
 - Hyperparameters-tuning could be further explored

- SMC steering and activation additions all had a few hyperparameters that could've been fine tuned given more time
 - better prompts for prompt steering could be investigated
 - using prompts like “you are a top-tier journalist at FOX/CNN news”
 - Ellie
- Conclusion (half a column)
 - Brief recap of results / interpretation
 - Implications for future studies
 - Future work include side effects of the methods
 - eg) we observe smc steering to produce worse quality generations and take much longer, but it would be nice to have quantitative evaluation of this
 - Rebecca